

C++ Programming II

STL - Concurrent Programming I

C++ Programming II
October 22, 2018

Prof. Dr. P. Arnold
Bern University of Applied Sciences

Agenda

► Intro

Lecture 5

Prof. Dr. P. Arnold



Bern University
of Applied Sciences

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Agenda

▶ Intro

▶ STL Threads



Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Agenda

- ▶ Intro
- ▶ STL Threads
- ▶ Data Races and Mutex



Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Agenda

- ▶ **Intro**
- ▶ **STL Threads**
- ▶ **Data Races and Mutex**
- ▶ **Mutexes**



Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Agenda

- ▶ **Intro**
- ▶ **STL Threads**
- ▶ **Data Races and Mutex**
- ▶ **Mutexes**
- ▶ **Locks**
 - ▶ Deadlock

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Agenda

- ▶ **Intro**
- ▶ **STL Threads**
- ▶ **Data Races and Mutex**
- ▶ **Mutexes**
- ▶ **Locks**
 - ▶ Deadlock
- ▶ **Call Once**



Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Intro

Intro

STL Threads


Data Races and Mutex

Mutexes

Locks


Deadlock

Call Once



- ▶ 1978: Intel 8086, 16-bit, 10MHz
- ▶ 2017: Intel Core i7-7700K, 64-bit, 4.2 GHz (4 cores (8 threads))

Stuttering



- ▶ Moore's Law: Double clock speed each 2 years
- ▶ But clock speed stabilized around 2000 (heat dissipation)
- ▶ Number of Cores did not!

Intro

STL Threads

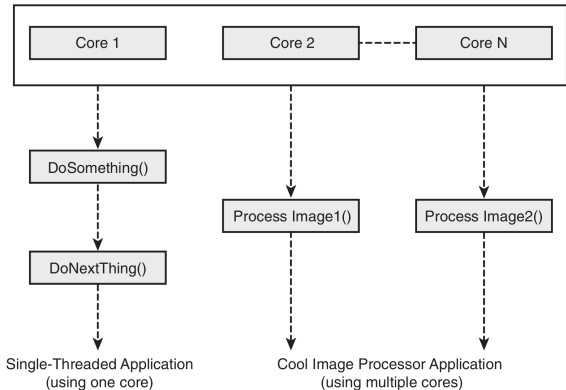
Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once



- ▶ Single Core: Sequential Program Flow
- ▶ Multi Core: Parallel Program Flow
- ▶ Since C++ 11, STL provides a clean and simple way to start and stop threads without using external libraries

STL Threads

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Enable Threads in CMake

`find_package` and `target_link_libraries`

- ▶ In order to use STL threads we have to:
 1. locate them with `find_package` and
 2. link the OS specific thread libraries with `target_link_libraries`

Enable Threads in CMake

find_package and target_link_libraries

- ▶ In order to use STL threads we have to:
 1. locate them with `find_package` and
 2. link the OS specific thread libraries with `target_link_libraries`
- ▶ Fortunately, CMake is doing that for us:

```
...  
...  
find_package(Threads)  
...  
...  
...  
target_link_libraries (${PROJECT_NAME} ${CMAKE_THREAD_LIBS_INIT})
```

- ▶ **Note:** Linking might not be necessary on Windows!

```
1  #include <iostream>
2
3  using namespace std;
4
5  void counting()
6  {
7      for (size_t i = 0; i < 1000; ++i)
8          cout << i << endl;
9  }
10
11 int main()
12 {
13     counting();
14     return 0;
15 }
```

- ▶ Single threaded program which counts to 1000

```
1 #include <iostream>
2 #include <thread>
3
4 using namespace std;
5
6 void counting()
7 {
8     for (size_t i = 0; i < 1000; ++i)
9         cout << i << endl;
10 }
11
12 int main()
13 {
14     thread t(counting);
15     t.join();
16
17     return 0;
18 }
```

- ▶ In the main thread we start a second thread doing the work.
- ▶ The main thread waits for the “worker-thread” to finish by calling the `join()` function! Otherwise the application might crash.
- ▶ Either the function `join()` or `detach()` have to be used.
- ▶ `join()`: blocks the main thread until the “worker-thread” is finished!
- ▶ `detach()`: detaches the “worker-thread” from the main thread, no longer external control over the thread


```
1 #include <iostream>
2 #include <thread>
3
4 using namespace std;
5
6 void counting(size_t count)
7 {
8     for (size_t i = 0; i < count; ++i)
9         cout << i << endl;
10 }
11
12 int main()
13 {
14     thread t(counting, 1000);
15     t.join();
16
17     return 0;
18 }
```

- We can also provide function parameters to the thread constructor

Thread Synchronisation

An Example

```
1 #include <iostream>
2 #include <thread>
3 using namespace std;
4 using namespace chrono_literals;
5
6 void threadWithParam(int threadNbr)
7 {
8     this_thread::sleep_for(1ms * threadNbr);
9     cout << "Hello from thread " << threadNbr << '\n';
10
11     this_thread::sleep_for(1s * threadNbr);
12     cout << "Bye from thread " << threadNbr << '\n';
13 }
```

- ▶ Simple function taking an argument as thread ID to identify the thread output later
- ▶ The threads wait a different amount of time to not write to `cout` at the same time

Thread Synchronisation

An Example

```
1 int main()
2 {
3     cout << thread::hardware_concurrency()
4         << " concurrent threads are supported.\n";
5     thread t1 {threadWithParam, 1};
6     thread t2 {threadWithParam, 2};
7     thread t3 {threadWithParam, 3};
8
9     t1.join();
10    t2.join();
11    t3.detach();
12
13    cout << "Threads joined.\n";
14 }
```

- ▶ In the main function we can print how many threads can be run at the same time using `hardware_concurrency()`
- ▶ We start 3 threads with different thread ID. Thread `t{f, x}` leads to a call of `f(x)`
- ▶ Since these threads are freely running, we need to stop them again when they are done with their work (`join()`, `detach()`)

Thread Synchronisation


An Example

- ▶ The output of the program can look like this:

```
8 concurrent threads are supported.  
Hello from thread 1  
Hello from thread 2  
Hello from thread 3  
Bye from thread 1  
Bye from thread 2  
Threads joined.
```

Thread Synchronisation

`join()` and `detach()`



- ▶ Reading the diagram from top to the bottom, it shows one point in time where we split the program workflow to four threads in total
- ▶ After starting the threads, the main thread executing the main function remained without work
- ▶ The main thread waits for thread 1 and thread 2, but not for detached thread 3

Thread Parameters

Passing Value by Reference using `std::ref`

- ▶ By default, parameters are passed by copy
- ▶ To pass by reference use `std::ref`

```
#include <iostream>
#include <thread>
using namespace std;

void addToVal(int& val, int addent)
{
    val += addent;
}

int main()
{
    int val = 23;
    // Pass parameter by reference
    thread t{addToVal, std::ref(val), 19};
    // Main and child thread share memory!

    if(t.joinable())
        t.join();

    cout << " Result is: " << val << endl; // 42
    return 0;
}
```



Data Races and Mutex

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Data Race Example

Two Threads try to set the same variable `x`

```
#include <iostream>
#include <thread>

using namespace std;

void setX(int& x, int value)
{
    x = value;
}

int main()
{
    int x = 1;
    thread setX1{setX, ref(x), 100};
    thread setX2{setX, ref(x), 200};

    setX1.join();
    setX2.join();

    cout << x << " ";
    return 0;
}
```

- ▶ What is the output?
- ▶ Demo...

Data Race Example

Two Threads try to set the same variable **x**

```
#include <iostream>
#include <thread>

using namespace std;

void setX(int& x, int value)
{
    x = value;
}

int main()
{
    int x = 1;
    thread setX1{setX, ref(x), 100};
    thread setX2{setX, ref(x), 200};

    setX1.join();
    setX2.join();

    cout << x << " ";
    return 0;
}
```

- ▶ What is the output?
- ▶ Demo...
- ▶ The two threads are racing for the same resource

Data Race Example 2

Two Threads race for the common resource `cout`

```
#include <iostream>
#include <string>
#include <thread>

using namespace std;

void print()
{
    for (int i = 0; i > -100; --i)
        cout << "From child: " << i << endl;
}

int main()
{
    thread t{print};

    for (int i = 0; i < 100; ++i)
        cout << "From main: " << i << endl;

    t.join();

    return 0;
}
```

- ▶ The output is cluttered!
- ▶ We have to synchronise the access of the common resource `cout`!

Data Race Example 2

Two Threads race for the common resource cout

```
From child: -29
From child: -30
From child: -31
From child: -32
From main: From child: -33
28From child: -34
From child: -35

From main: 29
From main: 30
From main: 31
From main: 32
From child: From main: -36
From child: -37
From child: -38
From child: -39
33From child: -40
From child: -41
```

Data Race Example 2

Synchronise the access of the common resource with `std::mutex`

```
1 #include <iostream>
2 #include <string>
3 #include <thread>
4 #include <mutex>
5
6 using namespace std;
7
8 mutex mu;
9
10 void sharedPrint(string msg, int value)
11 {
12     mu.lock(); // only one thread can enter!
13     cout << msg << value << endl;
14     mu.unlock();
15 }
```

- ▶ The term mutex stands for **mutual exclusion**
- ▶ We can restrict access to one thread using `std::mutex`
- ▶ `lock` locks the mutex
- ▶ `unlock` frees the mutex

Data Race Example 2

Synchronise the access of the common resource with `std::mutex`

```
1  #include <iostream>
2  #include <string>
3  #include <thread>
4  #include <mutex>
5
6  using namespace std;
7
8  mutex mu;
9
10 void sharedPrint(string msg, int value)
11 {
12     mu.lock(); // only one thread can enter!
13     cout << msg << value << endl;
14     mu.unlock();
15 }
16
17 void print()
18 {
19     for (int i = 0; i > -100; --i)
20         sharedPrint("From child: ", i);
21 }
22
23 int main()
24 {
25     thread t{print};
26
27     for (int i = 0; i < 100; ++i)
28         sharedPrint("From main: ", i);
29
30     t.join();
31
32     return 0;
33 }
```



- ▶ Do you see a problem? What happens if the code between the lock fails, *i.e.* throws an exception
- ▶ Mutex `mu` is locked for ever, called a **deadlock**



Mutex helper classes

Control std::mutex with lock_guard following RAI

```
1 #include <iostream>
2 #include <string>
3 #include <thread>
4 #include <mutex>
5
6 using namespace std;
7
8 mutex mu;
9
10 void sharedPrint(string msg, int value)
11 {
12     lock_guard<mutex> guard(mu); // RAI
13     cout << msg << value << endl;
14 }
```

- ▶ Don't use lock and unlock directly
- ▶ Use RAI technique: lock_guard
- ▶ Mutex is freed automatically when going out of scope
- ▶ Risk of deadlock reduced



Mutexes

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Overview Mutex Classes

Many different mutex classes in STL

Besides `mutex`, STL provides many different kinds of mutexes:

1. `mutex`:
provides `lock`, `unlock` and non-blocking `try_lock` method



Overview Mutex Classes

Many different mutex classes in STL

Besides `mutex`, STL provides many different kinds of mutexes:

1. `mutex`:
provides `lock`, `unlock` and non-blocking `try_lock` method
2. `timed_mutex`:
Same as `mutex`, but provides *timing out* methods `try_lock_for` and `try_lock_until`

Overview Mutex Classes

Many different mutex classes in STL

Besides `mutex`, STL provides many different kinds of mutexes:

1. `mutex`:
provides `lock`, `unlock` and non-blocking `try_lock` method
2. `timed_mutex`:
Same as `mutex`, but provides *timing out* methods `try_lock_for` and `try_lock_until`
3. `recursive_mutex`:
 - ▶ Same as `mutex`, but `mutex` can be locked multiple times by the same thread without blocking.
 - ▶ It is released after the owning thread called `unlock` as often as it called `lock`.

Overview Mutex Classes

Many different mutex classes in STL

Besides `mutex`, STL provides many different kinds of mutexes:

1. `mutex`:
provides `lock`, `unlock` and non-blocking `try_lock` method
2. `timed_mutex`:
Same as `mutex`, but provides *timing out* methods `try_lock_for` and `try_lock_until`
3. `recursive_mutex`:
 - ▶ Same as `mutex`, but `mutex` can be locked multiple times by the same thread without blocking.
 - ▶ It is released after the owning thread called `unlock` as often as it called `lock`.
4. `recursive_timed_mutex`:
Provides the features of both `timed_mutex` and `recursive_mutex`

Overview Mutex Classes

Many different mutex classes in STL

Besides `mutex`, STL provides many different kinds of mutexes:

1. `mutex`:
provides `lock`, `unlock` and non-blocking `try_lock` method
2. `timed_mutex`:
Same as `mutex`, but provides *timing out* methods `try_lock_for` and `try_lock_until`
3. `recursive_mutex`:
 - ▶ Same as `mutex`, but `mutex` can be locked multiple times by the same thread without blocking.
 - ▶ It is released after the owning thread called `unlock` as often as it called `lock`.
4. `recursive_timed_mutex`:
Provides the features of both `timed_mutex` and `recursive_mutex`
5. `shared_mutex`:
 - ▶ It can be locked in *exclusive* mode and in *shared* mode.
 - ▶ If a thread locks it in shared mode, it is possible for other threads to lock it in shared mode, too!
 - ▶ While a lock is locked in shared mode, it is not possible to obtain exclusive ownership.
 - ▶ This is very similar to the behavior of `shared_ptr`, only that it does not manage memory, but lock ownership.





Locks

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Overview RAII Lock-helpers

`std::lock_guard`

For memory management, we have `unique_ptr`, `shared_ptr` and `weak_ptr`. Those helpers provide very convenient ways to avoid memory leaks. Such helpers exist for mutexes, too.



Overview RAII Lock-helpers

`std::lock_guard`

For memory management, we have `unique_ptr`, `shared_ptr` and `weak_ptr`. Those helpers provide very convenient ways to avoid memory leaks. Such helpers exist for mutexes, too.

- ▶ C++ 11 provides two flavors of locks
 1. `std::lock_guard` for the simple use cases
 2. `std::unique_lock` for the advanced use cases

Overview RAII Lock-helpers

`std::lock_guard`

For memory management, we have `unique_ptr`, `shared_ptr` and `weak_ptr`. Those helpers provide very convenient ways to avoid memory leaks. Such helpers exist for mutexes, too.

- ▶ C++ 11 provides two flavors of locks
 1. `std::lock_guard` for the simple use cases
 2. `std::unique_lock` for the advanced use cases
- ▶ C++ 17 provides two more locks
 1. `std::scoped_lock` supporting multiple mutexes
 2. `std::shared_lock` locks mutex in shared mode (not discussed)

Overview RAII Lock-helpers

`std::lock_guard`

For memory management, we have `unique_ptr`, `shared_ptr` and `weak_ptr`. Those helpers provide very convenient ways to avoid memory leaks. Such helpers exist for mutexes, too.

- ▶ C++ 11 provides two flavors of locks
 1. `std::lock_guard` for the simple use cases
 2. `std::unique_lock` for the advanced use cases
- ▶ C++ 17 provides two more locks
 1. `std::scoped_lock` supporting multiple mutexes
 2. `std::shared_lock` locks mutex in shared mode (not discussed)

RAII Lock-helpers

`std::lock_guard`

- ▶ The simplest one is `std::lock_guard`

```
{  
    std::mutex m;  
    std::lock_guard<std::mutex> lockGuard{m};  
    sharedVariable= getVar();  
}
```

- ▶ `std::lock_guard` element's constructor accepts a mutex, on which it calls lock immediately.
- ▶ The whole constructor call will block until it obtains the lock on the mutex



RAII Lock-helpers

`std::unique_lock`

- ▶ The more advanced is `std::unique_lock`
- ▶ Allows fine-grained locking

```
std::mutex m;  
std::unique_lock<std::mutex> locker(m, std::defer_lock);  
// Locker owns mutex, but m is not locked
```

RAII Lock-helpers

`std::unique_lock`

- ▶ The more advanced is `std::unique_lock`
- ▶ Allows fine-grained locking

```
std::mutex m;  
std::unique_lock<std::mutex> locker(m, std::defer_lock);  
// Locker owns mutex, but m is not locked  
  
// Do something  
  
locker.lock(); // Now the mutex is locked  
// Work on shared resource
```

RAII Lock-helpers

`std::unique_lock`

- ▶ The more advanced is `std::unique_lock`
- ▶ Allows fine-grained locking


```
std::mutex m;  
std::unique_lock<std::mutex> locker(m, std::defer_lock);  
// Locker owns mutex, but m is not locked  
  
// Do something  
  
locker.lock(); // Now the mutex is locked  
// Work on shared resource  
  
locker.unlock(); // Now the mutex is unlocked again  
// Do something else
```

RAII Lock-helpers

`std::unique_lock`

- ▶ The more advanced is `std::unique_lock`
- ▶ Allows fine-grained locking

```
std::mutex m;  
std::unique_lock<std::mutex> locker(m, std::defer_lock);  
// Locker owns mutex, but m is not locked  
  
// Do something  
  
locker.lock(); // Now the mutex is locked  
// Work on shared resource  
  
locker.unlock(); // Now the mutex is unlocked again  
// Do something else  
  
locker.lock(); // Now the mutex is locked again  
// etc.
```



- ▶ If multiple mutexes are locked in the wrong sequence by multiple threads, a deadlock can happen
- ▶ Thread1 gets mutex1 and tries to get mutex2, Thread2 gets mutex2 and tries to get mutex1, which is locked by thread1 → deadlock

Deadlock

Avoiding deadlocks with `std::scoped_lock`

```
1 using namespace std;  
2 using namespace chrono_literals;  
3  
4 mutex mutA;  
5 mutex mutB;
```



Deadlock

Avoiding deadlocks with `std::scoped_lock`

```
1 using namespace std;
2 using namespace chrono_literals;
3
4 mutex mutA;
5 mutex mutB;
6
7 void deadlockFunc1()
8 {
9     cout << "bad f1 acquiring mutex A..." << endl;
10    lock_guard<mutex> la{mutA};
11    this_thread::sleep_for(100ms);
12    cout << "bad f1 acquiring mutex B..." << endl;
13    lock_guard<mutex> lb{mutB};
14    cout << "bad f1 got both mutexes." << endl;
15 }
16
17 void deadlockFunc2()
18 {
19     cout << "bad f2 acquiring mutex B..." << endl;
20    lock_guard<mutex> lb{mutB};
21    this_thread::sleep_for(100ms);
22    cout << "bad f2 acquiring mutex A..." << endl;
23    lock_guard<mutex> la{mutA};
24    cout << "bad f2 got both mutexes." << endl;
25 }
```





```
1 using namespace std;
2 using namespace chrono_literals;
3
4 mutex mutA;
5 mutex mutB;
6
7 void deadlockFunc1()
8 {
9     cout << "bad f1 acquiring mutex A..." << endl;
10    lock_guard<mutex> la{mutA};
11    this_thread::sleep_for(100ms);
12    cout << "bad f1 acquiring mutex B..." << endl;
13    lock_guard<mutex> lb{mutB};
14    cout << "bad f1 got both mutexes." << endl;
15 }
16
17 void deadlockFunc2()
18 {
19     cout << "bad f2 acquiring mutex B..." << endl;
20    lock_guard<mutex> lb{mutB};
21    this_thread::sleep_for(100ms);
22    cout << "bad f2 acquiring mutex A..." << endl;
23    lock_guard<mutex> la{mutA};
24    cout << "bad f2 got both mutexes." << endl;
25 }
```

- In main we create two threads:

```
1 thread t1{deadlockFunc1};
2 thread t2{deadlockFunc2};
3 t1.join();
4 t2.join();
```

- A typical deadlock - the program will hang forever!

- ▶ With `scoped_lock` STL provides a solution

```
1 void saneFunc1 ()
2 {
3     scoped_lock l{mutA, mutB};
4     cout << "sane f1 got both mutexes." << endl;
5 }
6 void saneFunc2 ()
7 {
8     scoped_lock l{mutB, mutA};
9     cout << "sane f2 got both mutexes." << endl;
10 }
```

```
1     thread t1{saneFunc1};
2     thread t2{saneFunc2};
3     t1.join();
4     t2.join();
```

- ▶ `scoped_lock` uses the `std::lock` function, which applies a special algorithm that performs a series of `try_lock` calls on all the mutexes provided, in order to prevent deadlocking.
- ▶ It's perfectly safe to use `scoped_lock` or call `std::lock` on the same set of locks, but in different orders.



Call Once

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once

Call Once Flag

If a task has to be done only once

- ▶ Imagine you want to add a header to your output only once, but you don't know which thread will be first
- ▶ We introduce a variable to make sure the header is only written once

```
1 bool hdrWritten{false};
2
3 void sharedPrint(string msg, int value)
4 {
5     if(!hdrWritten)
6     {
7         std::cout << std::string(4, '-')
8                     << " HEADER "
9                     << std::string(4, '-') << endl;
10
11         hdrWritten = true;
12     }
13
14     lock_guard<mutex> coutGuard(mu);
15     cout << msg << value << endl;
16 }
```

- ▶ This is not thread safe!
- ▶ We have to put the guard above!

Call Once Flag

If a task has to be done only once

- ▶ STL provides an elegant solution for exactly this kind of issue
- ▶ `once_flag` and `call_once`

```
1 once_flag flag;  
2 void sharedPrint(string msg, int value)  
3 {  
4     call_once(flag, []{std::cout << std::string(4, '-')  
5                          << " HEADER " << std::string(4, '-')  
6                          << endl;});  
7  
8     lock_guard<mutex> coutGuard(mu);  
9     cout << msg << value << endl;  
10 }
```

- ▶ This way the program is safe and efficient

Thank You

Questions

???

Lecture 5

Prof. Dr. P. Arnold



Bern University
of Applied Sciences

Intro

STL Threads

Data Races and Mutex

Mutexes

Locks

Deadlock

Call Once